

Code Review: a5-navya-shabbir-yu

Bhanu, Joyal, Jiangtao

October 22, 2017

Contents

Details	1
Report Review	1
Code Review	1
Execution	3
Final Thoughts	3

Details

Review

a5-navya-shabbir-yu

Reviewer

Bhanu, Joyal, Jiangtao

Codebase

<https://github.ccs.neu.edu/pdpmr-fl7/a5-navya-shabbir-yu>

Report Review

- Report is well structured and explains each component well.

Code Review

- Should have completion check on Job Result.

File: a5-navya-shabbir-yu/src/main/java/org/neu/pdpmr/tasks/Main.java

```
96:         new ExtractData(this.inDir, this.outDir, trainDatDir, qryYearDataDir, this.maxTrainYear
```

At L96 you can check whether `ExtractData` is completed or not. Any exception other than `IOException`, `InterruptedException`, `ClassNotFoundException` will be consumed at job level and there is no mechanism to know the job status other than the return value of `job.waitForCompletion`. If `job.waitForCompletion == -1`, you should halt any further execution. With current code, even if `job.waitForCompletion == -1` we still move on to the next task.

- Invalid Data Formats for Time Fields

```
File: a5-navya-shabbir-yu/src/main/java/org/neu/pdpmr/tasks/mappers/DataExtractingMapper.java
66:         float timeZone           = crsArrTime - crsDeptTime - crsElapsedTime;
67:         float cancelledCheck      = timeZone - actElapsedTime;

crsArrTime, crsDeptTime, crsElapsedTime considered as float while they are actually in time format
hhmm
```

- Avoid Consecutive Casting

```
File: a5-navya-shabbir-yu/src/main/java/org/neu/pdpmr/tasks/mappers/DataExtractingMapper.java
46:         context.write(new IntWritable((int) ((float) f.get("YEAR"))), datin);
47:     } catch (Exception e) {
```

Year is a `int`, but was taken as `float`. As a result authors had to cast from `Object` to `float` to `int`. This pattern is seen everywhere in the code. You can set `year` to `int` in your parser, as `year` can never be `float`

- Authors are parsing CSV records twice in both `DataExtractingMapper` and `DataExtractingReducer`. It is unnecessary compute. You can parse once in `DataExtractingMapper` and extract only meaningful fields. These fields can be converted into custom writable values and emitted from mapper. Thus in the reducer you don't have to parse CSV texts again.
- You can get the below year directly from the key.

```
File: a5-navya-shabbir-yu/src/main/java/org/neu/pdpmr/tasks/reducers/DataExtractingReducer.java
63:         // year in the record
64:         int recordYear = (int) ((float) f.get("YEAR"));
65:
```

- `DataExtractingReducer.LOC101`: `min` is already 1, so `max - min + 1` would always be `max`

```
File: a5-navya-shabbir-yu/src/main/java/org/neu/pdpmr/tasks/reducers/DataExtractingReducer.java
098:         Random rand = new Random(System.currentTimeMillis());
099:         int max = 10000;
100:         int min = 1;
101:         int res = rand.nextInt(max - min + 1) + min;
```

- Instead of manually merging `MultiPartOutput` you can use `LazyOutputFormat.setOutputFormatClass(job, TextOutputFormat.class)` to get merged results from reducer. (Assuming you need to merge only the multipart output.)
- Unnecessary Query serialization/serialization

Reading queries in `Main.java` -> converting it to Query Collection -> converting it to JSON string -> setting this string into job config -> converting this string to Query Collection in `Job2`

This seems to be a very complicated process with unnecessary conversions and parsing. You can simply read query file in `job2`.

- The authors have used counters, but there is no explanation of what these counters are doing. Cannot understand just by reading the code.

```
File: a5-navya-shabbir-yu/src/main/java/org/neu/pdpmr/tasks/mappers/TwoHopFlightsMapper.java
118:         // flightDate >= journeyDateMs && flightDate <= journeyDateMs + X_DAYS
119:         context.getCounter("MyMap", "s1").increment(1);
120:         if (!(flightDateMs >= q.dateMs)) return;
121:         context.getCounter("MyMap", "s2").increment(1);
122:         if (!(flightDateMs <= q.dateMs + THREE_DAYS)) return;
123:         context.getCounter("MyMap", "s3").increment(1);
124:
```

- The join operation to generate two hop flights is very complex and missing documentation. Cannot comprehend just by reading the code.

Execution

- The implementation has memory issues, and the code does not run completely as mentioned in the report.
- `make all` is missing `exec` target to run hadoop job.

Final Thoughts

- Cannot verify the accuracy as the results are not generated.
- Code quality is not very good and hence making it very difficult to read.
- Documentation about joins and code structure would be very helpful.
- Random Forest is generally a memory intensive algorithm. You can try simple models such as `average by year`, or linear classifications to optimize memory usage.
- You can avoid using complex data structures such as `TypeToken` and `GSON` serializations.